

# final

April 29, 2024

```
[6]: import tensorflow as tf
      from tensorflow.keras.datasets import fashion_mnist
      from sklearn.model_selection import train_test_split
```

```
[8]: # Load the Fashion MNIST dataset
      (train_images, train_labels), (test_images, test_labels) = fashion_mnist.
      ↪load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>

29515/29515 0s 1us/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>

26421880/26421880 1s

0us/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>

5148/5148 0s 1us/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>

4422102/4422102 0s

0us/step

```
[9]: # Normalize pixel values to be in the range [0, 1]
      train_images = train_images / 255.0
      test_images = test_images / 255.0
```

```
[10]: # Split the dataset into training and testing sets
       train_images, val_images, train_labels, val_labels = ↪
       ↪train_test_split(train_images, train_labels, test_size=0.2, random_state=42)
```

```
[11]: from tensorflow.keras import layers, models
```

```
[13]: # Define the input shape
       input_shape = (28, 28, 1)

       # Define the CNN architecture
       model_cnn = models.Sequential([
```

```

layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
layers.MaxPooling2D((2, 2)),
layers.Conv2D(64, (3, 3), activation='relu'),
layers.MaxPooling2D((2, 2)),
layers.Conv2D(64, (3, 3), activation='relu'),
layers.Flatten(),
layers.Dense(64, activation='relu'),
layers.Dense(10, activation='softmax')
])

```

```

[14]: # Compile the model
model_cnn.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

```

```

[15]: # Reshape the input data for CNN (add a channel dimension)
train_images_cnn = train_images.reshape((train_images.shape[0], 28, 28, 1))
val_images_cnn = val_images.reshape((val_images.shape[0], 28, 28, 1))

```

```

[16]: # Train the model
history_cnn = model_cnn.fit(train_images_cnn, train_labels, epochs=10,
                             ↪batch_size=64,
                             validation_data=(val_images_cnn, val_labels))

```

```

Epoch 1/10
750/750      8s 10ms/step -
accuracy: 0.7150 - loss: 0.8106 - val_accuracy: 0.8495 - val_loss: 0.4085
Epoch 2/10
750/750      8s 10ms/step -
accuracy: 0.8661 - loss: 0.3720 - val_accuracy: 0.8771 - val_loss: 0.3309
Epoch 3/10
750/750      7s 10ms/step -
accuracy: 0.8866 - loss: 0.3136 - val_accuracy: 0.8898 - val_loss: 0.2982
Epoch 4/10
750/750      8s 10ms/step -
accuracy: 0.9010 - loss: 0.2743 - val_accuracy: 0.8942 - val_loss: 0.2911
Epoch 5/10
750/750      7s 10ms/step -
accuracy: 0.9103 - loss: 0.2486 - val_accuracy: 0.8987 - val_loss: 0.2772
Epoch 6/10
750/750      7s 10ms/step -
accuracy: 0.9184 - loss: 0.2233 - val_accuracy: 0.8951 - val_loss: 0.2862
Epoch 7/10
750/750      7s 10ms/step -
accuracy: 0.9258 - loss: 0.2065 - val_accuracy: 0.9053 - val_loss: 0.2580
Epoch 8/10
750/750      7s 10ms/step -
accuracy: 0.9311 - loss: 0.1852 - val_accuracy: 0.9038 - val_loss: 0.2620

```

```
Epoch 9/10
750/750          7s 10ms/step -
accuracy: 0.9364 - loss: 0.1707 - val_accuracy: 0.9043 - val_loss: 0.2645
Epoch 10/10
750/750          7s 10ms/step -
accuracy: 0.9409 - loss: 0.1578 - val_accuracy: 0.9120 - val_loss: 0.2542
```

```
[17]: # Evaluate the model on validation data
val_loss, val_acc = model_cnn.evaluate(val_images_cnn, val_labels)
print("Validation accuracy:", val_acc)
```

```
375/375          1s 2ms/step -
accuracy: 0.9140 - loss: 0.2482
Validation accuracy: 0.9120000004768372
```

```
[18]: # Reshape the test data
test_images_cnn = test_images.reshape((test_images.shape[0], 28, 28, 1))
```

```
[19]: # Evaluate the model on test data
test_loss, test_acc = model_cnn.evaluate(test_images_cnn, test_labels)
print("Test accuracy:", test_acc)
```

```
313/313          1s 2ms/step -
accuracy: 0.9018 - loss: 0.2820
Test accuracy: 0.9039999842643738
```

```
[21]: # Define the DNN architecture
model_dnn = models.Sequential([
    layers.Input(shape=(28, 28)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(10, activation='softmax')
])
```

```
[22]: # Compile the DNN model
model_dnn.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```

```
[23]: # Train the DNN model
history_dnn = model_dnn.fit(train_images, train_labels, epochs=10,
                             ↪batch_size=64,
                             validation_data=(val_images, val_labels))
```

```
Epoch 1/10
750/750          1s 1ms/step -
```

```

accuracy: 0.7023 - loss: 0.8469 - val_accuracy: 0.8462 - val_loss: 0.4300
Epoch 2/10
750/750          1s 943us/step -
accuracy: 0.8388 - loss: 0.4513 - val_accuracy: 0.8595 - val_loss: 0.3907
Epoch 3/10
750/750          1s 906us/step -
accuracy: 0.8546 - loss: 0.4076 - val_accuracy: 0.8637 - val_loss: 0.3669
Epoch 4/10
750/750          1s 1ms/step -
accuracy: 0.8618 - loss: 0.3808 - val_accuracy: 0.8697 - val_loss: 0.3576
Epoch 5/10
750/750          1s 997us/step -
accuracy: 0.8709 - loss: 0.3575 - val_accuracy: 0.8750 - val_loss: 0.3463
Epoch 6/10
750/750          1s 982us/step -
accuracy: 0.8726 - loss: 0.3464 - val_accuracy: 0.8777 - val_loss: 0.3302
Epoch 7/10
750/750          1s 934us/step -
accuracy: 0.8798 - loss: 0.3332 - val_accuracy: 0.8765 - val_loss: 0.3459
Epoch 8/10
750/750          1s 889us/step -
accuracy: 0.8815 - loss: 0.3265 - val_accuracy: 0.8778 - val_loss: 0.3247
Epoch 9/10
750/750          1s 905us/step -
accuracy: 0.8875 - loss: 0.3075 - val_accuracy: 0.8813 - val_loss: 0.3237
Epoch 10/10
750/750          1s 916us/step -
accuracy: 0.8877 - loss: 0.3005 - val_accuracy: 0.8838 - val_loss: 0.3202

```

```

[24]: # Evaluate the DNN model on validation data
      val_loss_dnn, val_acc_dnn = model_dnn.evaluate(val_images, val_labels)
      print("Validation accuracy (DNN):", val_acc_dnn)

```

```

375/375          0s 281us/step -
accuracy: 0.8862 - loss: 0.3210
Validation accuracy (DNN): 0.8837500214576721

```

```

[25]: # Evaluate the DNN model on test data
      test_loss_dnn, test_acc_dnn = model_dnn.evaluate(test_images, test_labels)
      print("Test accuracy (DNN):", test_acc_dnn)

```

```

313/313          0s 281us/step -
accuracy: 0.8793 - loss: 0.3455
Test accuracy (DNN): 0.8777999877929688

```

```

[26]: # Define the MLP architecture
      model_mlp = models.Sequential([
          layers.Flatten(input_shape=(28, 28)),

```

```

layers.Dense(256, activation='relu'),
layers.Dropout(0.3),
layers.Dense(128, activation='relu'),
layers.Dropout(0.3),
layers.Dense(64, activation='relu'),
layers.Dropout(0.3),
layers.Dense(10, activation='softmax')
])

```

```

[27]: # Compile the MLP model
model_mlp.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

```

```

[28]: # Train the MLP model
history_mlp = model_mlp.fit(train_images, train_labels, epochs=10,
                             ↪batch_size=64,
                             validation_data=(val_images, val_labels))

```

```

Epoch 1/10
750/750          2s 2ms/step -
accuracy: 0.6327 - loss: 1.0226 - val_accuracy: 0.8369 - val_loss: 0.4456
Epoch 2/10
750/750          1s 2ms/step -
accuracy: 0.8260 - loss: 0.4989 - val_accuracy: 0.8544 - val_loss: 0.3912
Epoch 3/10
750/750          1s 2ms/step -
accuracy: 0.8438 - loss: 0.4419 - val_accuracy: 0.8673 - val_loss: 0.3727
Epoch 4/10
750/750          1s 2ms/step -
accuracy: 0.8552 - loss: 0.4119 - val_accuracy: 0.8682 - val_loss: 0.3602
Epoch 5/10
750/750          1s 2ms/step -
accuracy: 0.8601 - loss: 0.3947 - val_accuracy: 0.8641 - val_loss: 0.3765
Epoch 6/10
750/750          1s 2ms/step -
accuracy: 0.8655 - loss: 0.3769 - val_accuracy: 0.8716 - val_loss: 0.3507
Epoch 7/10
750/750          1s 2ms/step -
accuracy: 0.8703 - loss: 0.3643 - val_accuracy: 0.8735 - val_loss: 0.3400
Epoch 8/10
750/750          1s 2ms/step -
accuracy: 0.8678 - loss: 0.3582 - val_accuracy: 0.8760 - val_loss: 0.3327
Epoch 9/10
750/750          1s 2ms/step -
accuracy: 0.8761 - loss: 0.3427 - val_accuracy: 0.8820 - val_loss: 0.3284
Epoch 10/10
750/750          1s 2ms/step -

```

accuracy: 0.8761 - loss: 0.3432 - val\_accuracy: 0.8808 - val\_loss: 0.3278

```
[29]: # Evaluate the MLP model on validation data
val_loss_mlp, val_acc_mlp = model_mlp.evaluate(val_images, val_labels)
print("Validation accuracy (MLP):", val_acc_mlp)
```

375/375                      0s 429us/step -  
accuracy: 0.8815 - loss: 0.3267  
Validation accuracy (MLP): 0.8807500004768372

```
[30]: # Evaluate the MLP model on test data
test_loss_mlp, test_acc_mlp = model_mlp.evaluate(test_images, test_labels)
print("Test accuracy (MLP):", test_acc_mlp)
```

313/313                      0s 409us/step -  
accuracy: 0.8741 - loss: 0.3505  
Test accuracy (MLP): 0.8748999834060669

```
[31]: #As committee
import numpy as np
```

```
[32]: # Get the predicted probabilities for each model on the test data
probabilities_cnn = model_cnn.predict(test_images)
probabilities_dnn = model_dnn.predict(test_images)
probabilities_mlp = model_mlp.predict(test_images)
```

313/313                      1s 2ms/step  
313/313                      0s 283us/step  
313/313                      0s 374us/step

```
[33]: # Average the predicted probabilities across all three models
ensemble_probabilities = (probabilities_cnn + probabilities_dnn +  
    probabilities_mlp) / 3
```

```
[34]: # Get the class predictions by selecting the class with the highest probability
ensemble_predictions = np.argmax(ensemble_probabilities, axis=1)
```

```
[35]: # Evaluate the ensemble on the test data
test_accuracy_ensemble = np.mean(ensemble_predictions == test_labels)
print("Test accuracy (Ensemble):", test_accuracy_ensemble)
```

Test accuracy (Ensemble): 0.9045

```
[37]: # Report individual model accuracies
print("Individual Model Accuracies:")
print("CNN Model Accuracy:", test_acc)
print("DNN Model Accuracy:", test_acc_dnn)
print("MLP Model Accuracy:", test_acc_mlp)

# Report final accuracy of the committee
```

```
print("\nFinal Accuracy of the Committee:")
print("Ensemble Model Accuracy:", test_accuracy_ensemble)
```

Individual Model Accuracies:

CNN Model Accuracy: 0.9039999842643738

DNN Model Accuracy: 0.8777999877929688

MLP Model Accuracy: 0.8748999834060669

Final Accuracy of the Committee:

Ensemble Model Accuracy: 0.9045

[ ]: *#Q2A How are the diverse deep-learning models formed in your project, please explain? (10points)*

```
[73]: print("CNN (Convolutional Neural Network):")
print("CNNs tend to be very good in capturing spatial hierarchies in images. In my implementation of CNNs it consists of convolutional layers followed by max-pooling layers and dense layers. The convolutional layers extract certain features from the images through filters.\n ")
print("DNN (Deep Neural Networks):")
print("DNNs consist of connected layers and are effective when learning patterns that are complex in data. In my implementation it is a simple feed forward network with multiple hidden layers. Each layer connects with another layer, this allows it to learn relations that are non-linear.\n ")
print("MLP (Multi-Layer perceptron):")
print("They are feedforward networks with multiple layers of nodes, they can be used for both classification and regression. In my implementation it has multiple hidden layers with ReLU activation functions and a softmax output layer for classification.\n")
```

CNN (Convolutional Neural Network):

CNNs tend to be very good in capturing spatial hierarchies in images. In my implementation of CNNs it consists of convolutional layers followed by max-pooling layers and dense layers. The convolutional layers extract certain features from the images through filters.

DNN (Deep Neural Networks):

DNNs consist of connected layers and are effective when learning patterns that are complex in data. In my implementation it is a simple feed forward network with multiple hidden layers. Each layer connects with another layer, this allows it to learn relations that are non-linear.

MLP (Multi-Layer perceptron):

They are feedforward networks with multiple layers of nodes, they can be used for both classification and regression. In my implementation it has multiple hidden layers with ReLU activation functions and a softmax output layer for classification.

```
[ ]: #2B How are the three diverse deep-learning models combined? (20 points).  
print("Done on Submission PDF")
```

```
[66]: #Q2C  
from sklearn.metrics import accuracy_score, precision_score, recall_score,  
      ↪ f1_score  
  
# Calculate metrics for individual models  
metrics_individual = {  
    "Model": ["CNN", "DNN", "MLP"],  
    "Accuracy": [test_acc, test_acc_dnn, test_acc_mlp]  
}  
  
# Calculate precision, recall, and F1-score for each individual model  
for model, model_name in zip([model_cnn, model_dnn, model_mlp], ["CNN", "DNN",  
      ↪ "MLP"]):  
    predictions = np.argmax(model.predict(test_images), axis=1)  
    metrics_individual[model_name + " Precision"] =  
    ↪ precision_score(test_labels, predictions, average='weighted')  
    metrics_individual[model_name + " Recall"] = recall_score(test_labels,  
    ↪ predictions, average='weighted')  
    metrics_individual[model_name + " F1-Score"] = f1_score(test_labels,  
    ↪ predictions, average='weighted')  
  
# Calculate metrics for the committee  
ensemble_predictions = np.argmax(ensemble_probabilities, axis=1)  
metrics_committee = {  
    "Model": ["Ensemble"],  
    "Accuracy": [test_accuracy_ensemble],  
    "Precision": [precision_score(test_labels, ensemble_predictions,  
    ↪ average='weighted')],  
    "Recall": [recall_score(test_labels, ensemble_predictions,  
    ↪ average='weighted')],  
    "F1-Score": [f1_score(test_labels, ensemble_predictions,  
    ↪ average='weighted')]  
}  
  
print("\nFinal Metrics of Individual Models:")  
print(pd.DataFrame(metrics_individual))  
  
print("\nFinal Metrics of the Committee:")  
print(pd.DataFrame(metrics_committee))
```

```
313/313          1s 2ms/step  
313/313          0s 265us/step  
313/313          0s 399us/step
```



Final Metrics of Individual Models:

	Model	Accuracy	CNN Precision	CNN Recall	CNN F1-Score	DNN Precision	\
0	CNN	0.9905	0.906548	0.904	0.90468	0.877663	
1	DNN	0.8778	0.906548	0.904	0.90468	0.877663	
2	MLP	0.8749	0.906548	0.904	0.90468	0.877663	

	DNN Recall	DNN F1-Score	MLP Precision	MLP Recall	MLP F1-Score
0	0.8778	0.875762	0.876199	0.8749	0.875076
1	0.8778	0.875762	0.876199	0.8749	0.875076
2	0.8778	0.875762	0.876199	0.8749	0.875076

Final Metrics of the Committee:

	Model	Accuracy	Precision	Recall	F1-Score
0	Ensemble	0.9045	0.904322	0.9045	0.904302

```
[ ]: print("Proper tabulated answers on pdf")
```